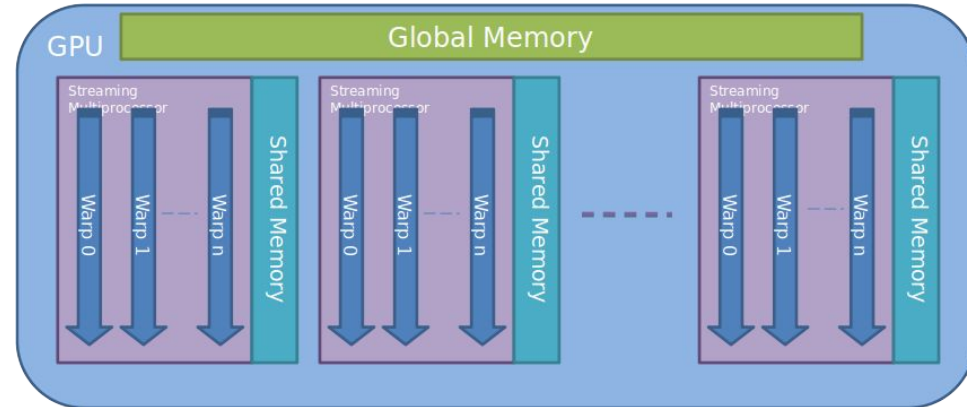


Optimizing Graph Algorithms for GPUs

By: Mayuresh Anand, Alon Albalak, Koa Sato

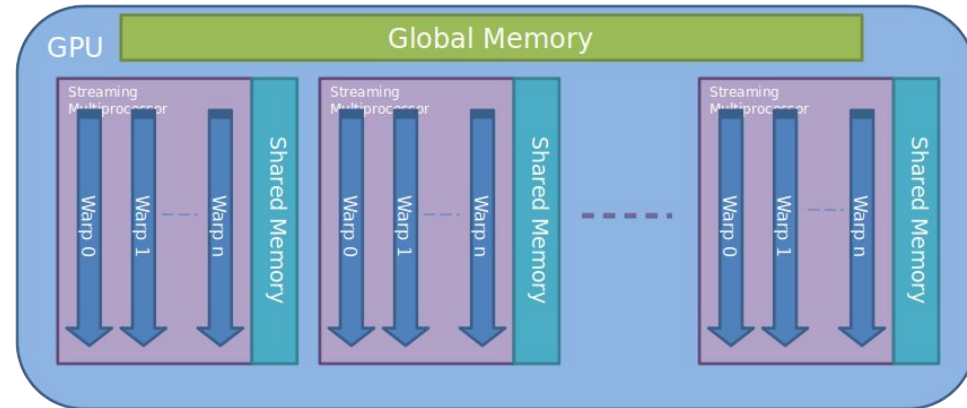
Graphical Processing Units

- Composed of GPU cores



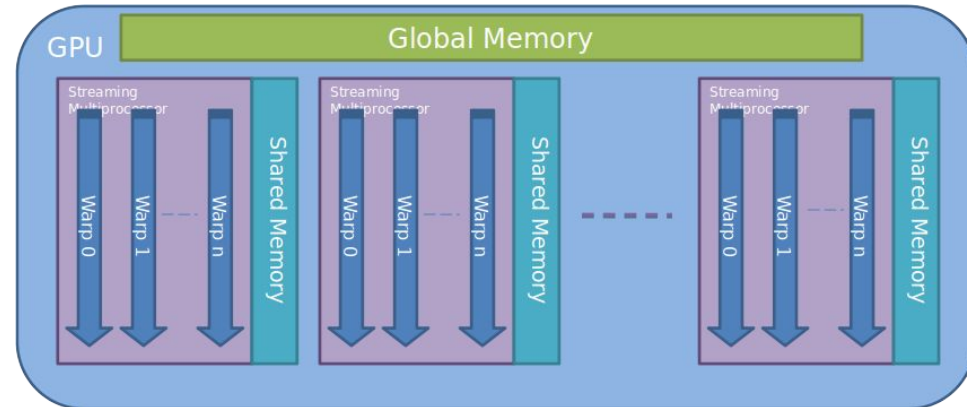
Graphical Processing Units

- Composed of GPU cores
- **Each core is composed of threads which are organized by a streaming multiprocessor (SM)**



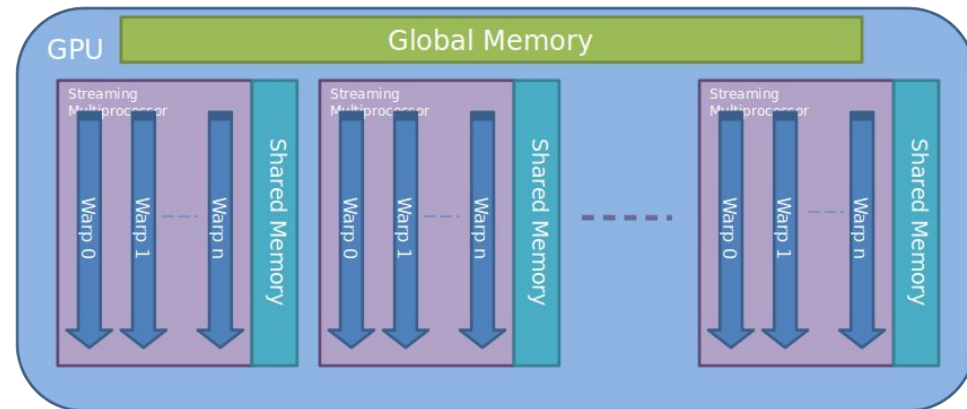
Graphical Processing Units

- Composed of GPU cores
- Each core is composed of threads which are organized by a streaming multiprocessor (SM)
- **Threads are organized into warps, such that each warp executes an instruction in parallel**



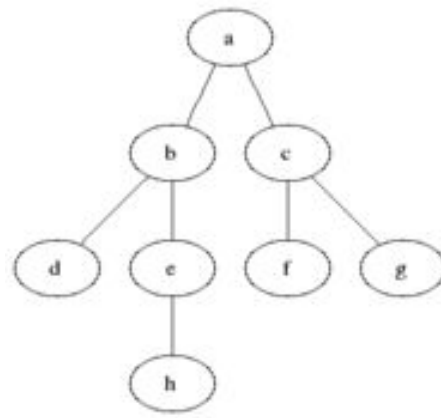
Graphical Processing Units

- Composed of GPU cores
- Each core is composed of threads which are organized by a streaming multiprocessor (SM)
- Threads are organized into warps, such that each warp executes an instruction in parallel
- **Most GPUs utilize single-instruction, multiple-threads paradigm**



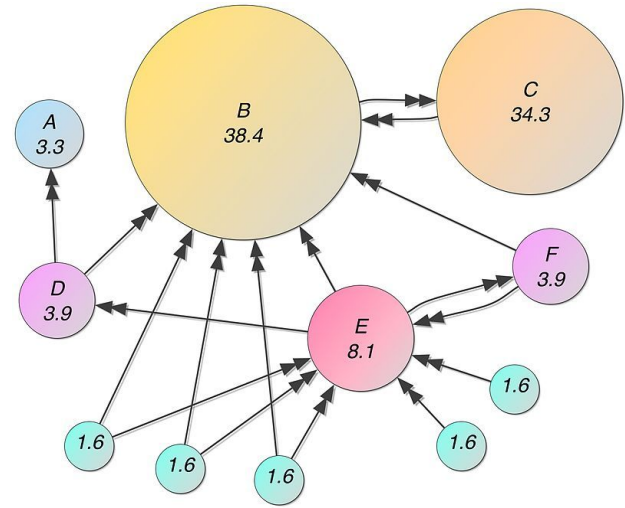
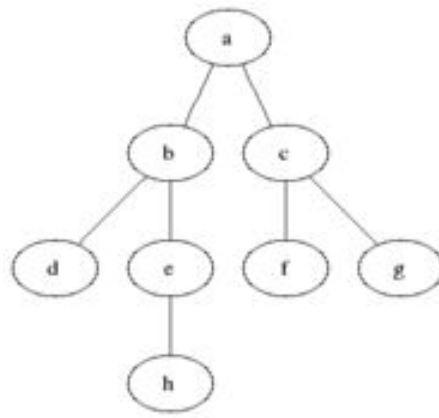
Graph Algorithms

- **Breadth-first search**



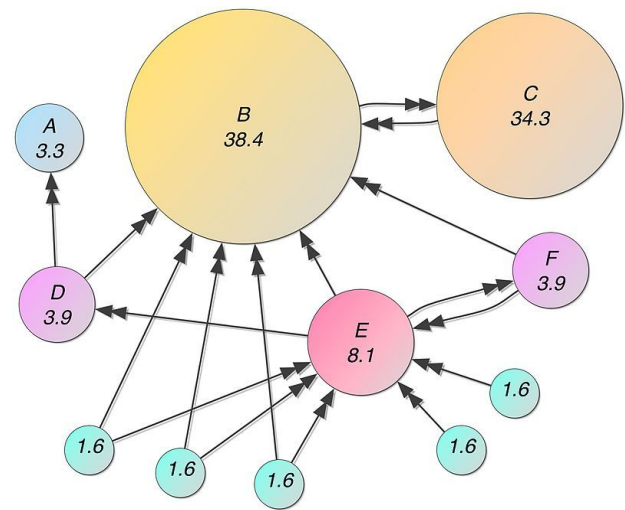
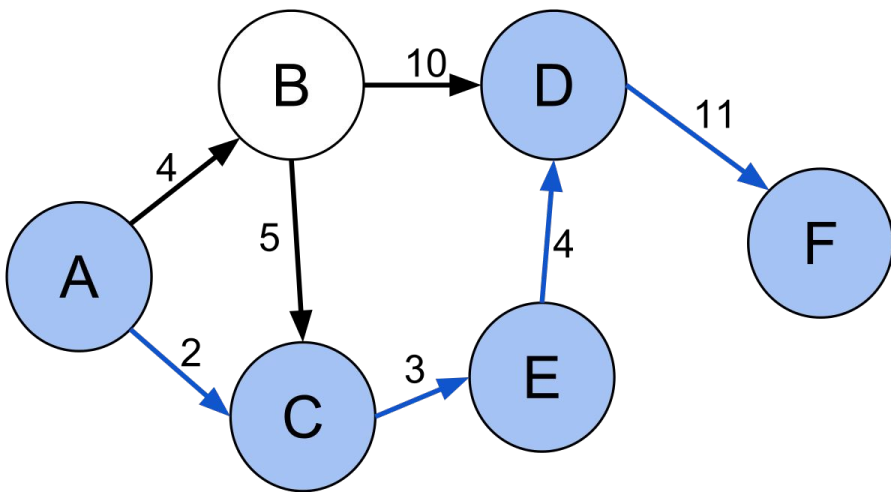
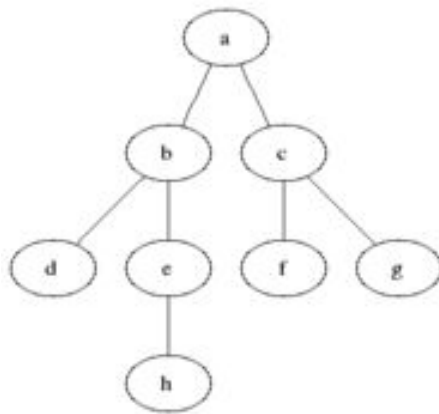
Graph Algorithms

- Breadth-first search
- **PageRank**

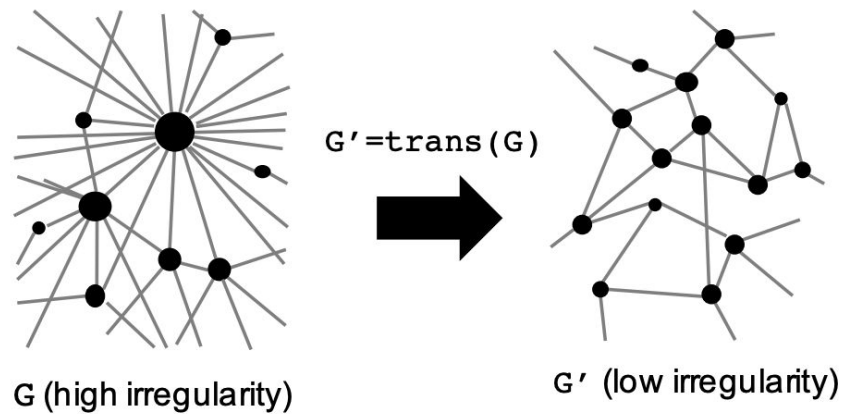


Graph Algorithms

- Breadth-first search
- PageRank
- **Single-source shortest path**

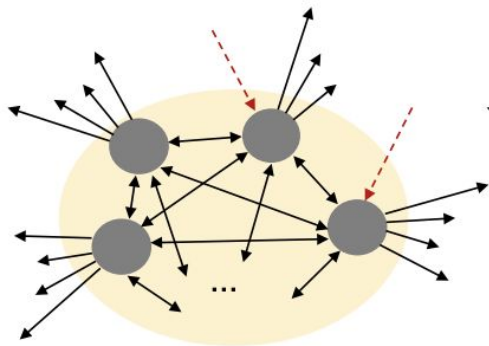


Graph Regularization

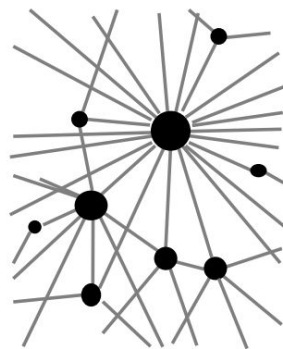


Graph Regularization

- Methods include transformations:
 - **Clique**

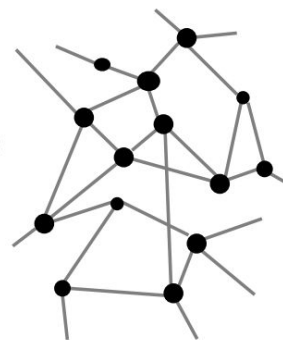
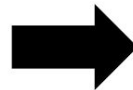


(a) \mathcal{T}_{cliq}



G (high irregularity)

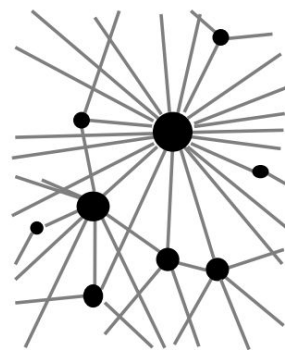
$$G' = \text{trans}(G)$$



G' (low irregularity)

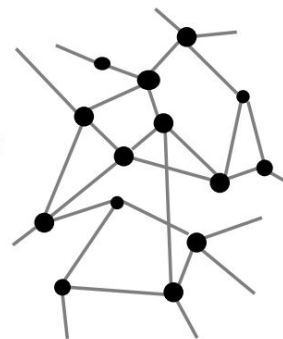
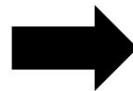
Graph Regularization

- Methods include transformations:
 - Clique
 - **Circular**

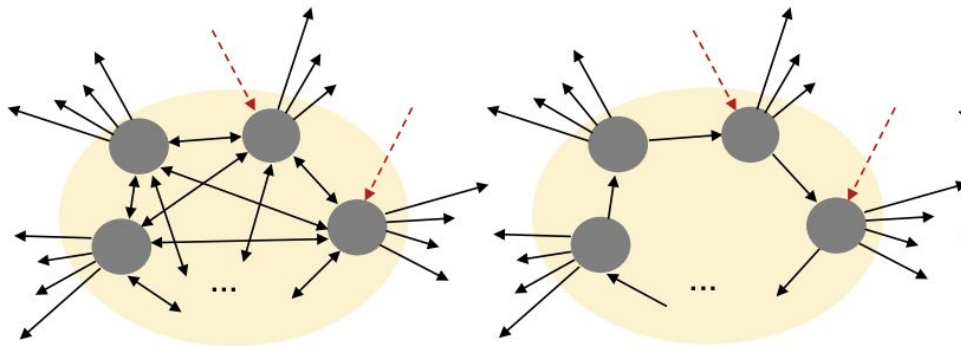


G (high irregularity)

$$G' = \text{trans}(G)$$



G' (low irregularity)

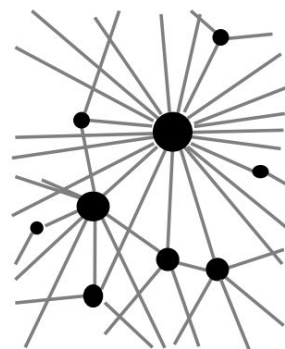


(a) \mathcal{T}_{cliq}

(b) \mathcal{T}_{circ}

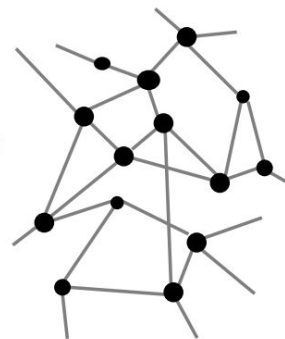
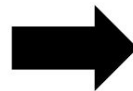
Graph Regularization

- Methods include transformations:
 - Clique
 - Circular
 - **Star**

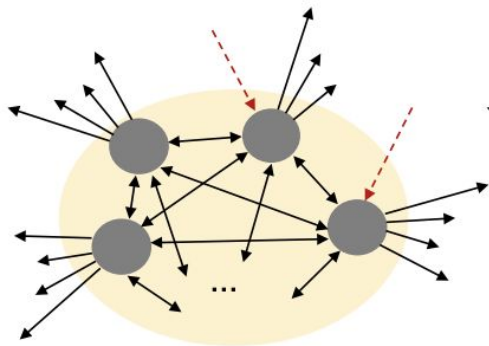


G (high irregularity)

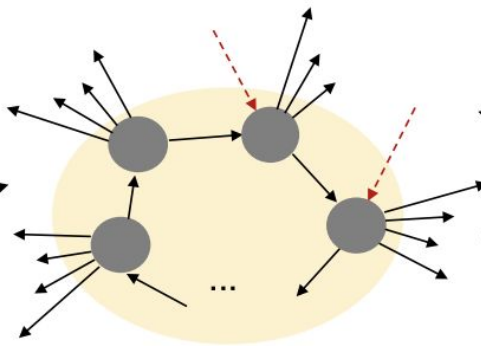
$$G' = \text{trans}(G)$$



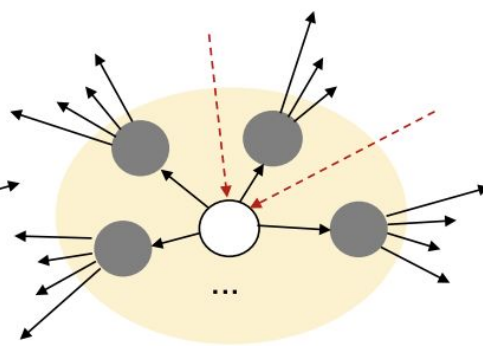
G' (low irregularity)



(a) \mathcal{T}_{cliq}



(b) \mathcal{T}_{circ}



(c) \mathcal{T}_{star}

Tigr - Transforming Irregular Graphs into more Regular ones

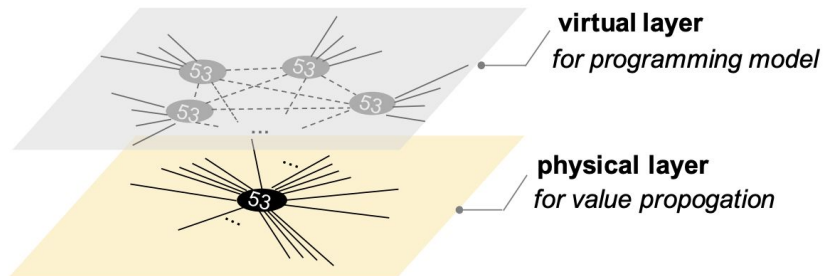
- Claim that GPU-based graph algorithms suffer due to power law structure

Tigr - Transforming Irregular Graphs into more Regular ones

- Claim that GPU-based graph algorithms suffer due to power law structure
- **Previous works addressed the issue by:**
 - Warp segmentation
 - Altering the graph algorithms

Tigr - Transforming Irregular Graphs into more Regular ones

- Claim that GPU-based graph algorithms suffer due to power law structure
- Previous works addressed the issue by:
 - Warp segmentation
 - Altering the graph algorithms
- **Tigr targets the fundamental issue of irregularity by transforming the graph *virtually***



CuSha

- Use shards!!!

CuSha

- Use shards!!!
- **Representing graph with shards has improved I/O performance for disk based graph processing**

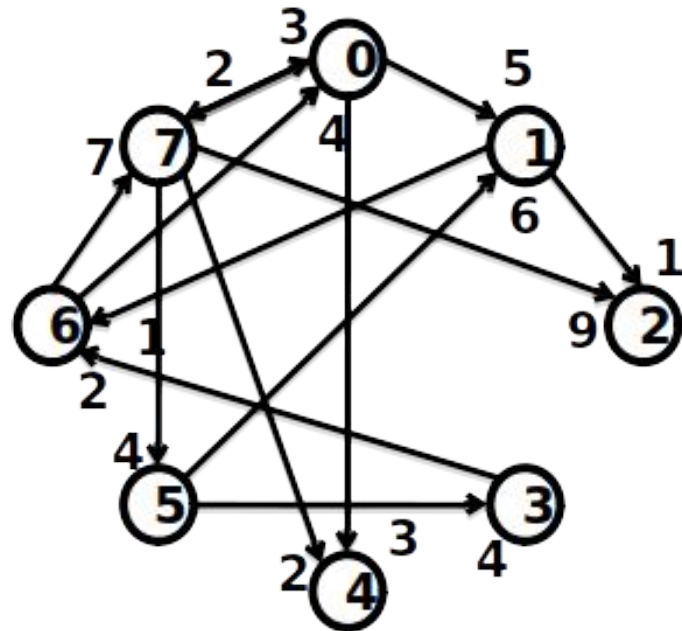
CuSha

- Use shards!!!
- Representing graph with shards has improved I/O performance for disk based graph processing
- **Shards allow contiguous placement of graph data required by a set of computation.**

CuSha

- Use shards!!!
- Representing graph with shards has improved I/O performance for disk based graph processing
- Shards allow contiguous placement of graph data required by a set of computation.
- **G-Shards representation employs shard concept and modifies it appropriately for GPU.**

CuSha



Shard 0

SrcIndex	SrcValue _e	EdgeValue	DestIndex
0	X ₀	5	1
1	X ₁	1	2
5	X ₅	6	1
5	X ₅	4	3
6	X ₆	4	0
7	X ₇	3	0
7	X ₇	9	2

Shard 1

SrcIndex	SrcValue _e	EdgeValue	DestIndex
0	X ₀	3	4
0	X ₀	2	7
1	X ₁	1	6
3	X ₃	2	6
6	X ₆	7	7
7	X ₇	2	4
7	X ₇	4	5

VertexValue_s

X	X	X	X	X	X	X	X
0	1	2	3	4	5	6	7

Motivation

- **We want to make graph algorithms run faster**

Motivation

- We want to make graph algorithms run faster
- **CUSHA uses G-shards to appropriately manage GPU resources**

Motivation

- We want to make graph algorithms run faster
- CUSHA uses G-shards to appropriately manage GPU resources
- **What if we convert graph to more regularized one using an intermediate representation and input to CUSHA**

Motivation

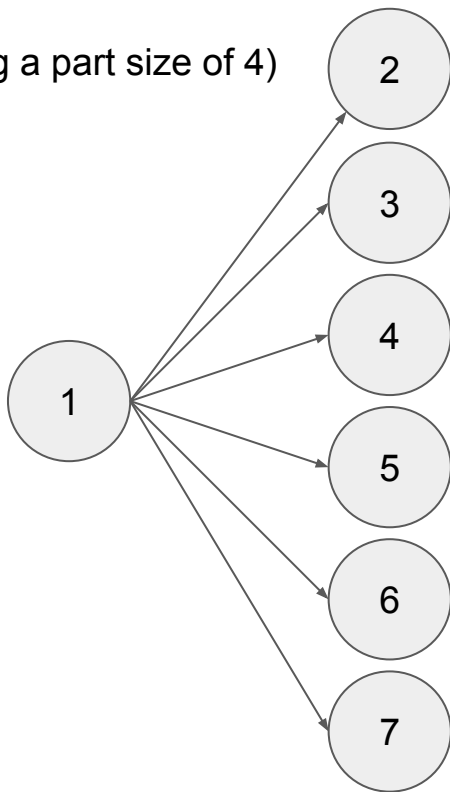
- We want to make graph algorithms run faster
- CUSHA uses G-shards to appropriately manage GPU resources
- What if we convert graph to more regularized one using an intermediate representation and input to CUSHA
- **This representation has lower out degree hence making it larger in size but regularized and efficient to process**

Hypothesis

- Can we achieve better performance from CuSha if we perform a physical graph transformation (UDT) that regularizes the graph before inputting it into CuSha?

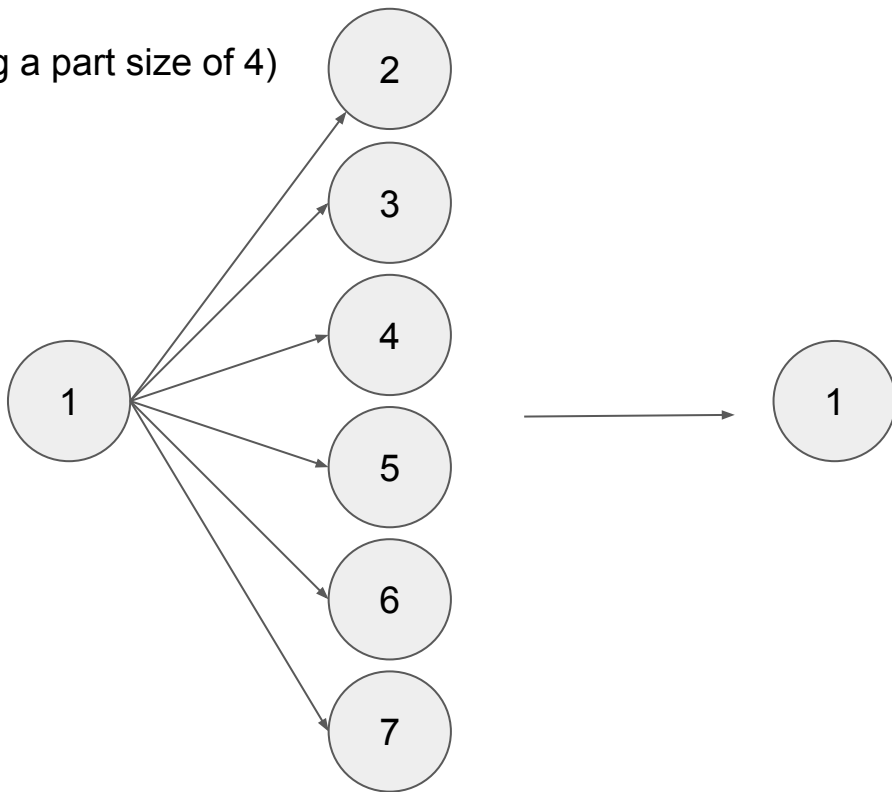
Uniform-Degree Tree Transformation

(Using a part size of 4)



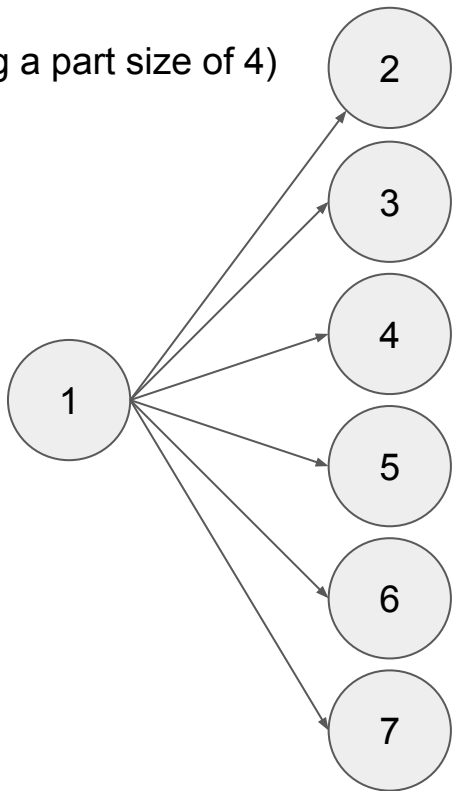
Uniform-Degree Tree Transformation

(Using a part size of 4)

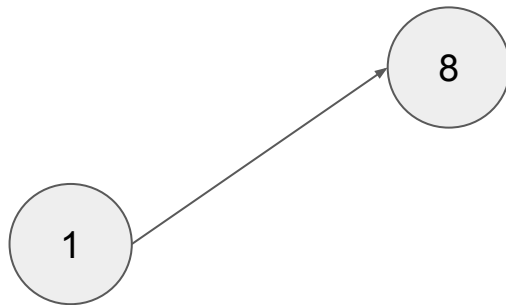


Uniform-Degree Tree Transformation

(Using a part size of 4)

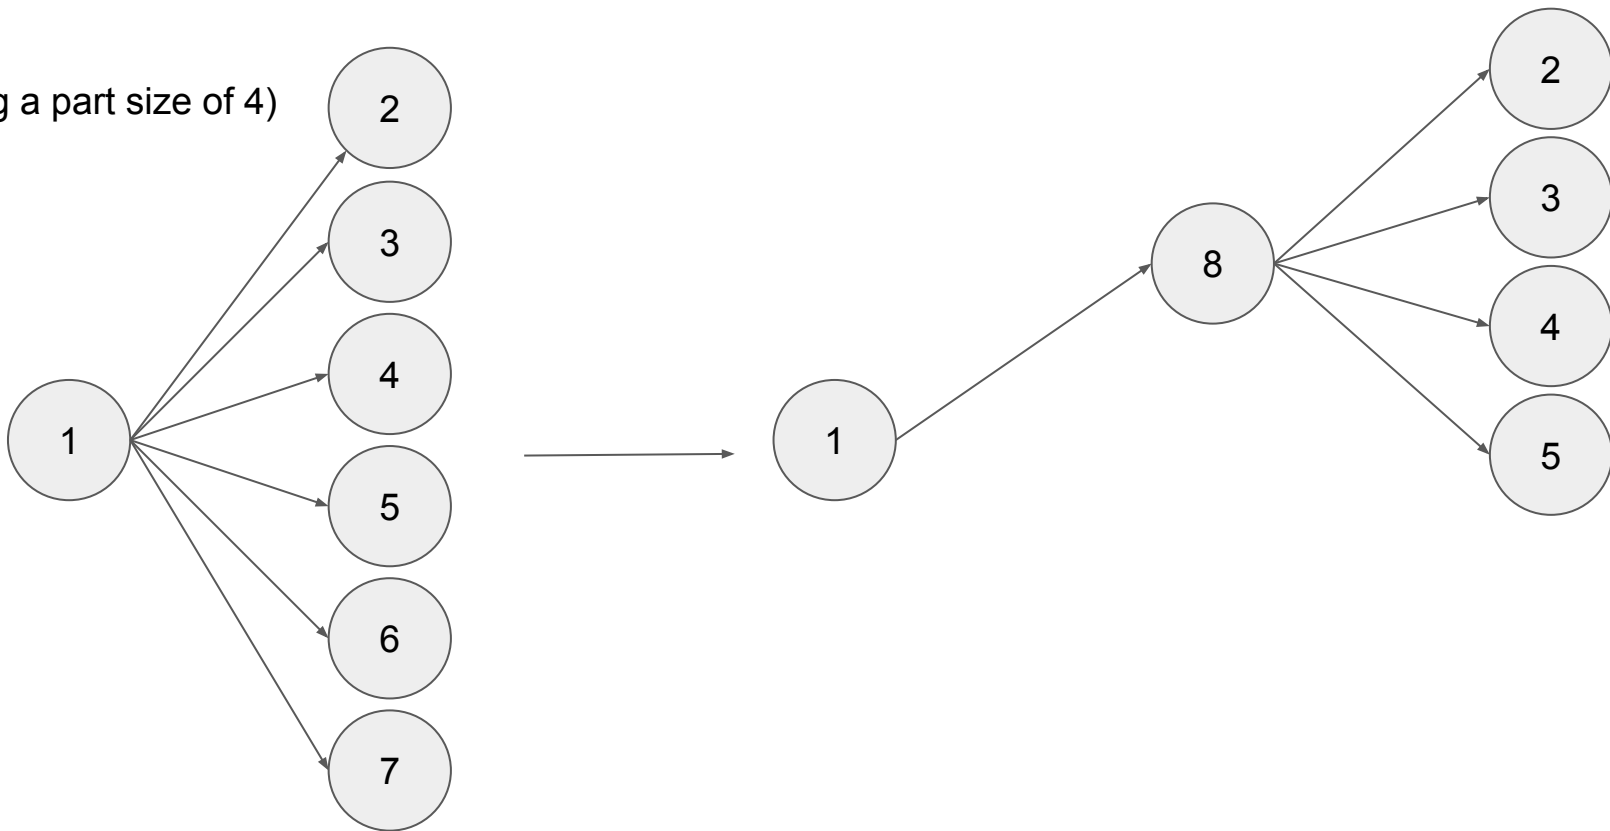


Add a new *virtual* node



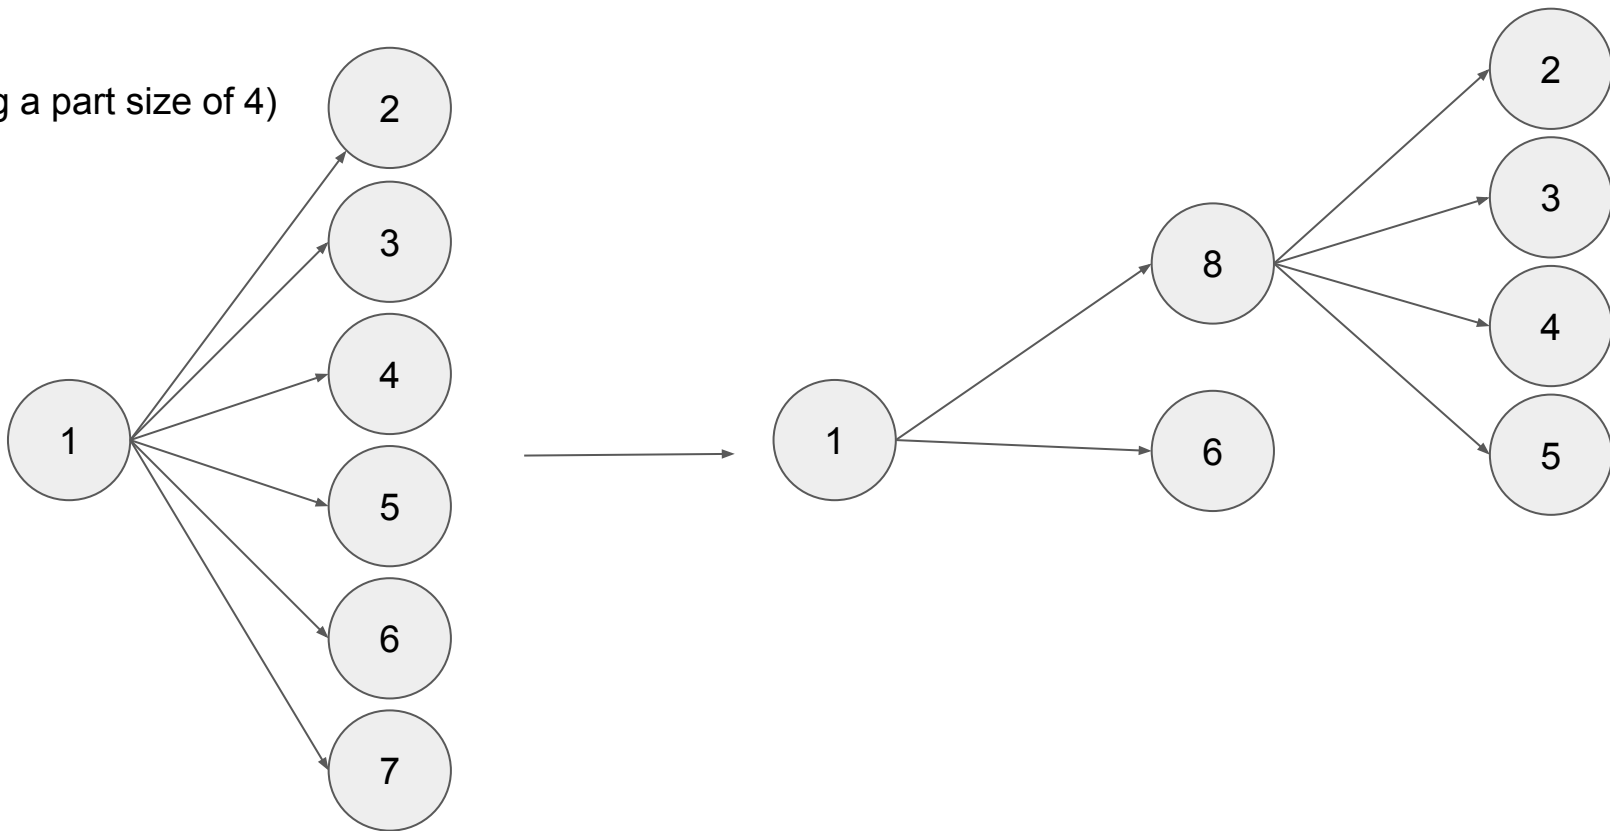
Uniform-Degree Tree Transformation

(Using a part size of 4)



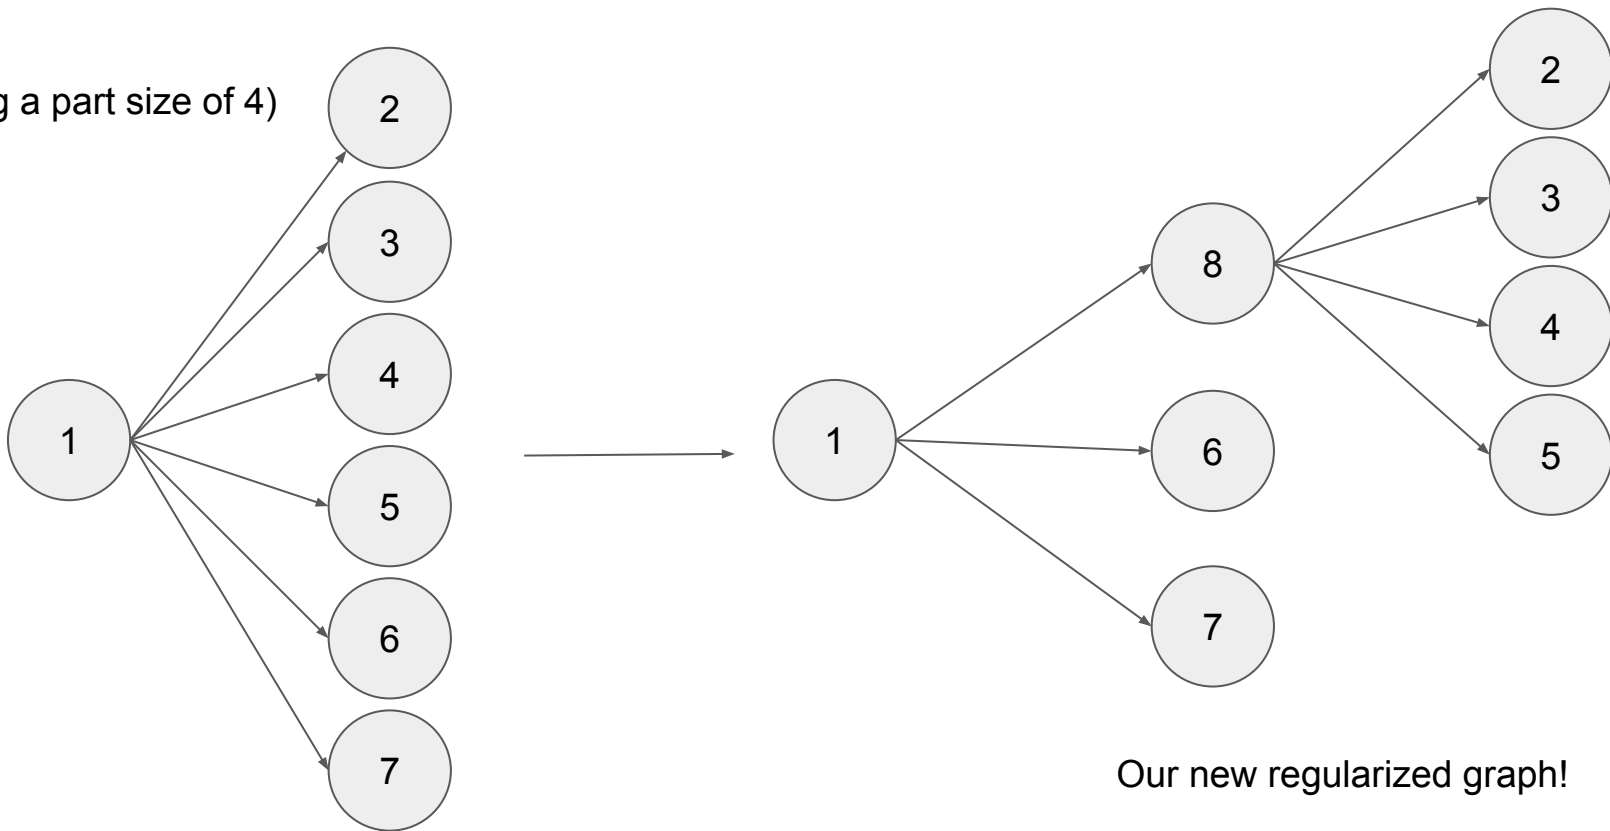
Uniform-Degree Tree Transformation

(Using a part size of 4)



Uniform-Degree Tree Transformation

(Using a part size of 4)



MACHINE CONFIGURATION

- CPU: i5 6600k (3.50 GHz, 4 cores, 4 threads)

MACHINE CONFIGURATION

- CPU: i5 6600k (3.50 GHz, 4 cores, 4 threads)
- **GPU: GTX 1060 6GB**

MACHINE CONFIGURATION

- CPU: i5 6600k (3.50 GHz, 4 cores, 4 threads)
- GPU: GTX 1060 6GB
- **Memory: 16 GB**

MACHINE CONFIGURATION

- CPU: i5 6600k (3.50 GHz, 4 cores, 4 threads)
- GPU: GTX 1060 6GB
- Memory: 16 GB
- **OS: Ubuntu 18.04**

Datasets

- Pokec social network¹
 - 1632803 vertices, 30622564 edges
- Higgs twitter²
 - 456627 vertices, 14855842 edges
- Amazon product co-purchasing network³
 - 403394 vertices, 3387388 edges

1 - <https://snap.stanford.edu/data/soc-Pokec.html>

2 - <https://snap.stanford.edu/data/higgs-twitter.html>

3 - <https://snap.stanford.edu/data/amazon0601.html>

EXPERIMENTATION

- **Computation time for UDT-CuSha on BFS and SSSP**

EXPERIMENTATION

- Computation time for UDT-CuSha on BFS and SSSP
- **Computation time on Higgs, Pokec and Amazon**

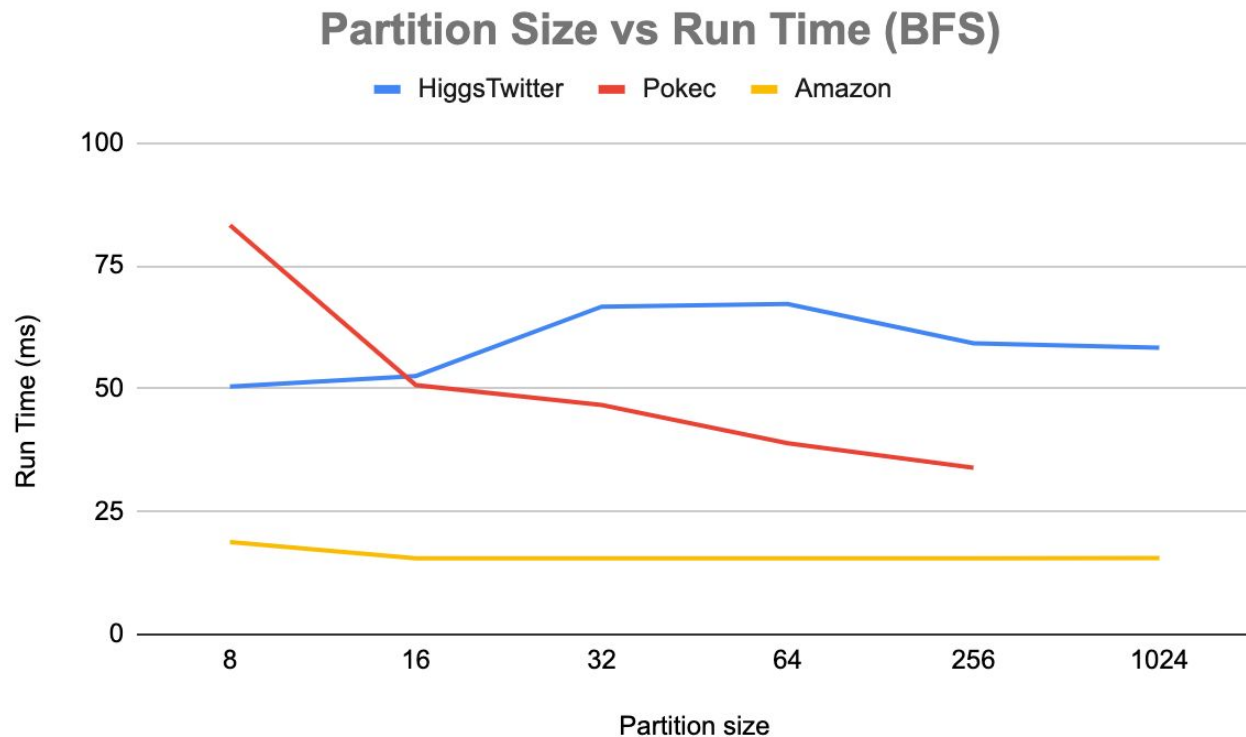
EXPERIMENTATION

- Computation time for UDT-CuSha on BFS and SSSP
- Computation time on Higgs, Pokec and Amazon
- **Computation time on 5 different part sizes**

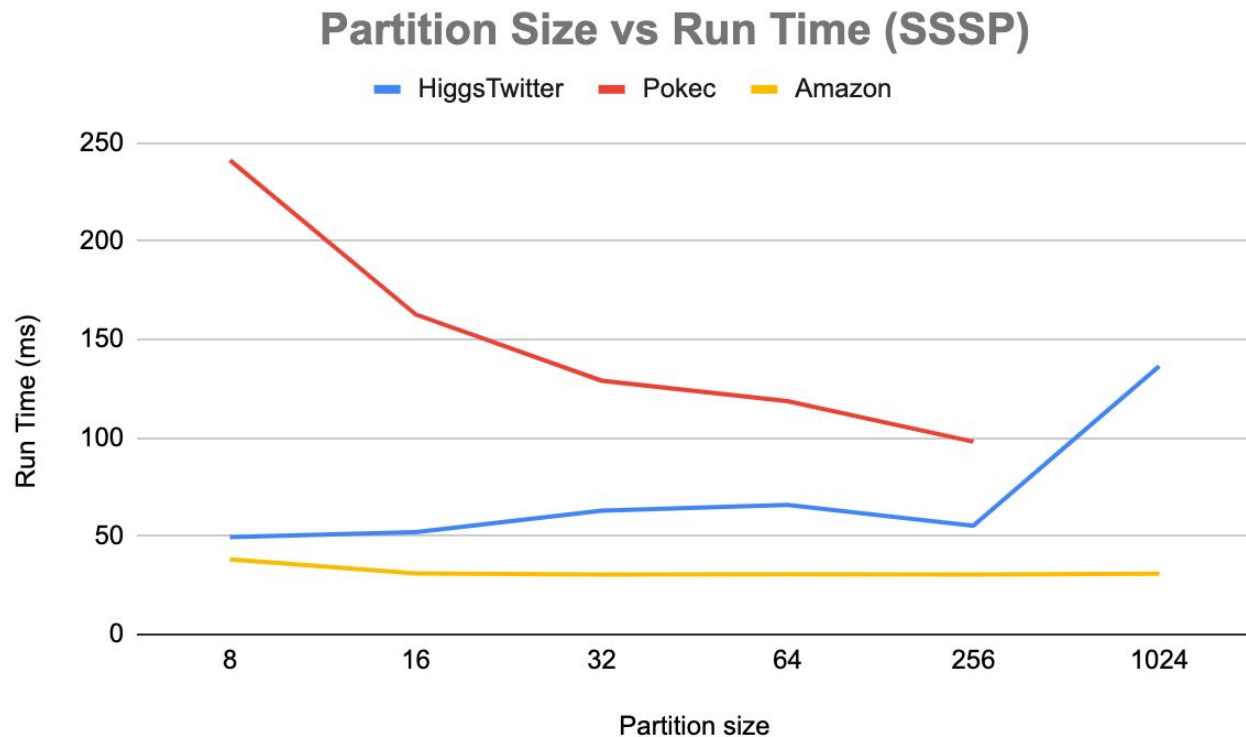
EXPERIMENTATION

- Computation time for UDT-CuSha on BFS and SSSP
- Computation time on Higgs, Pokec and Amazon
- Computation time on 5 different part sizes
- **Comparison with baseline (Tigr-V+ and CuSha)**

Experimentation Results



Experimentation Results



Comparisons with Baselines

Twitter	UDT-CuSha	UDT-CuSha	UDT-CuSha	UDT-CuSha	UDT-CuSha	UDT-CuSha	CuSha	Tigr
Part Size	8	16	32	64	256	1024	N/A	8
Time (ms)	49.4730	52.0758	62.9978	65.9644	55.3284	136.479	132.789	6.7826

Amazon	UDT-CuSha	UDT-CuSha	UDT-CuSha	UDT-CuSha	UDT-CuSha	UDT-CuSha	CuSha	Tigr
Part Size	8	16	32	64	256	1024	N/A	8
Time (ms)	38.1308	30.966	30.4466	30.5708	30.4798	30.8584	30.5586	7.182

Interpretations

- **Increasing part_size asymptotically leads to CuSha level performance**
 - Intuitive as increasing part size means we have done less to alter the graph

Interpretations

- Increasing `part_size` asymptotically leads to CuSha level performance
 - Intuitive as increasing part size means we have done less to alter the graph
- **Higgs vs. Amazon**
 - **As `part_size` increases, performance on Higgs decreases, performance on Amazon**
 - **Similar size graphs, however:**
 - **Higgs has 32.5 edges/node,**
 - **Amazon has 4.7 edges/node**
 - **Our approach works better on more dense graphs, which means regularization helps**

Interpretations

- Increasing part_size asymptotically leads to CuSha level performance
 - Intuitive as increasing part size means we have done less to alter the graph
- Higgs vs. Amazon
 - As part_size increases, performance on Higgs decreases, performance on Amazon
 - Similar size graphs, however:
 - Higgs has 32.5 edges/node,
 - Amazon has 4.7 edges/node
 - Our approach works better on more dense graphs, which means regularization helps
- **Tigr has incredible performance**
 - **Prevents physical transformation of graph while allowing computation on a normalized graph**

Interpretations

- Increasing part_size asymptotically leads to CuSha level performance
 - Intuitive as increasing part size means we have done less to alter the graph
- Higgs vs. Amazon
 - As part_size increases, performance on Higgs decreases, performance on Amazon
 - Similar size graphs, however:
 - Higgs has 32.5 edges/node,
 - Amazon has 4.7 edges/node
 - Our approach works better on more dense graphs, which means regularization helps
- Tigr has incredible performance
 - Prevents physical transformation of graph while allowing computation on a normalized graph
- **Trade-off between increased number of nodes and regularization**

Future work and things we didn't have time for

- **UDT transformation is useful, but only for graphs with high irregularity**

Future work and things we didn't have time for

- UDT transformation is useful, but only for graphs with high irregularity
- **Analysis of graph prior to execution of algorithms is very important**
 - A graph that has small maximum outgoing edges will not benefit from UDT

Future work and things we didn't have time for

- UDT transformation is useful, but only for graphs with high irregularity
- Analysis of graph prior to execution of algorithms is very important
 - A graph that has small maximum outgoing edges will not benefit from UDT
- **Try adjusting virtual warp size in CuSha algorithm to match part size**

QUESTIONS?